

# FinalProjectCST383

February 24, 2023

## 0.1 Project Pacific Analytics

### 0.1.1 David DeBow

### 0.1.2 Justin Johnson

### 0.1.3 Ryan Parker

### 0.1.4 Feb 15, 2023

Project progress report

Purpose. The purpose of this assignment is to help keep you on track, and to allow us to help you make adjustments to your project before it is due.

Exploration of Names

```
[1]: %matplotlib inline
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from matplotlib import rcParams
import seaborn as sns
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from sklearn.preprocessing import StandardScaler
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score
from sklearn.metrics import classification_report
from sklearn.model_selection import cross_val_score
import graphviz
# cite https://stackoverflow.com/questions/74877602/
# getting-a-warning-when-using-sklearn-neighbors-about-keepdims
from warnings import simplefilter
simplefilter(action='ignore', category=FutureWarning)

sns.set()
rcParams['figure.figsize'] = 8,6
sns.set_context('talk') # 'talk' for slightly larger
```

```
[2]: # code in this cell from:
# https://stackoverflow.com/questions/27934885/
# how-to-hide-code-from-cells-in-ipython-notebook-visualized-with-nbviewer
from IPython.display import HTML

HTML('''<script>
code_show=true;
function code_toggle() {
  if (code_show){
    $('div.input').hide();
  } else {
    $('div.input').show();
  }
  code_show = !code_show
}
$( document ).ready(code_toggle);
</script>
<form action="javascript:code_toggle()"><input type="submit" value="Click here_
to display/hide the code."></form>''')
```

[2]: <IPython.core.display.HTML object>

```
[3]: # The original CSV is too large for github. It is 31mb and github accepts 25mb_
# or less.
# The plan is to split the files and merge them on read.
# We found a website that can split CSVs online. https://www.splitcsv.com/
# Successfully split but ran in to problems
# For now we will host on Dropbox. Dropbox can handle larger files when not_
# using a webbrowser.
# Error, need to change domain to direct download. Went from dropbox.com to dl.
# dropboxusercontent.com
df = pd.read_csv("http://dl.dropboxusercontent.com/s/0uzku9gb353n3nv/
NYC_Dog_Licensing_Dataset.csv")

dfBaby = pd.read_csv("https://archive.ics.uci.edu/ml/machine-learning-databases/
00591/name_gender_dataset.csv")
```

The choice of dataset is Animal (Dog) License Data from NYC from data.gov . The file is approximately 31 megabytes and holds 500,000 lines of data. It is hosted on a dropbox for ease of use.

We are going to predict the AnimalName. There are 8 columns of data, but we are most interested in the AnimalName.

The predictors we will be using are ZipCode and Birthyear.

Preprocessing, exploration, and visualization, and machine learning work are below.

```
[4]: #Let us look at the information and general data of the chart.
df.info()
print(df.describe())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 508196 entries, 0 to 508195
Data columns (total 8 columns):
#   Column                Non-Null Count  Dtype
---  -
0   AnimalName            507169 non-null object
1   AnimalGender          508175 non-null object
2   AnimalBirthYear       508196 non-null int64
3   BreedName            508196 non-null object
4   ZipCode              508187 non-null float64
5   LicenseIssuedDate     508196 non-null object
6   LicenseExpiredDate    508119 non-null object
7   Extract Year          508196 non-null int64
dtypes: float64(1), int64(2), object(5)
memory usage: 31.0+ MB
```

	AnimalBirthYear	ZipCode	Extract Year
count	508196.000000	508187.000000	508196.000000
mean	2013.206304	10704.532395	2019.380810
std	4.847849	1096.178521	2.646618
min	1991.000000	0.000000	2016.000000
25%	2010.000000	10031.000000	2017.000000
50%	2014.000000	10468.000000	2018.000000
75%	2017.000000	11228.000000	2022.000000
max	2021.000000	99508.000000	2022.000000

```
data['Gender'].replace(0, 'Female',inplace=True) data['Gender'].replace(1, 'Male',inplace=True)
```

```
[5]: #Let us look at the information and general data of the chart.
dfBaby.info()
print(dfBaby.describe())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 147269 entries, 0 to 147268
Data columns (total 4 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Name                147269 non-null object
1   Gender              147269 non-null object
2   Count               147269 non-null int64
3   Probability         147269 non-null float64
dtypes: float64(1), int64(1), object(2)
memory usage: 4.5+ MB
```

	Count	Probability
count	1.472690e+05	1.472690e+05

mean	2.481161e+03	6.790295e-06
std	4.645472e+04	1.271345e-04
min	1.000000e+00	2.736740e-09
25%	5.000000e+00	1.368370e-08
50%	1.700000e+01	4.652460e-08
75%	1.320000e+02	3.612500e-07
max	5.304407e+06	1.451679e-02

```
[6]: #sns.pairplot(df)
      #df.plot()
```

```
[ ]:
```

How many Zipcodes are there?

```
[7]: df['ZipCode'].value_counts().head(10)
```

```
[7]: 10025.0    11439
      10023.0     9164
      11215.0     8829
      10024.0     8826
      11201.0     8757
      10011.0     8555
      10128.0     8392
      10009.0     8073
      10314.0     7613
      10312.0     7467
      Name: ZipCode, dtype: int64
```


```
[8]: df['ZipCode'].value_counts().size
```

```
[8]: 784
```

There are 740 zip codes.

What are the most popular animal names?

```
[9]: #@ 1 Show the top 20 names
      names = df['AnimalName'].value_counts().head(20)
      names.plot.bar()
      plt.xlabel("Name")
      plt.ylabel("Number")
      plt.title("Number of Names Animal");
```



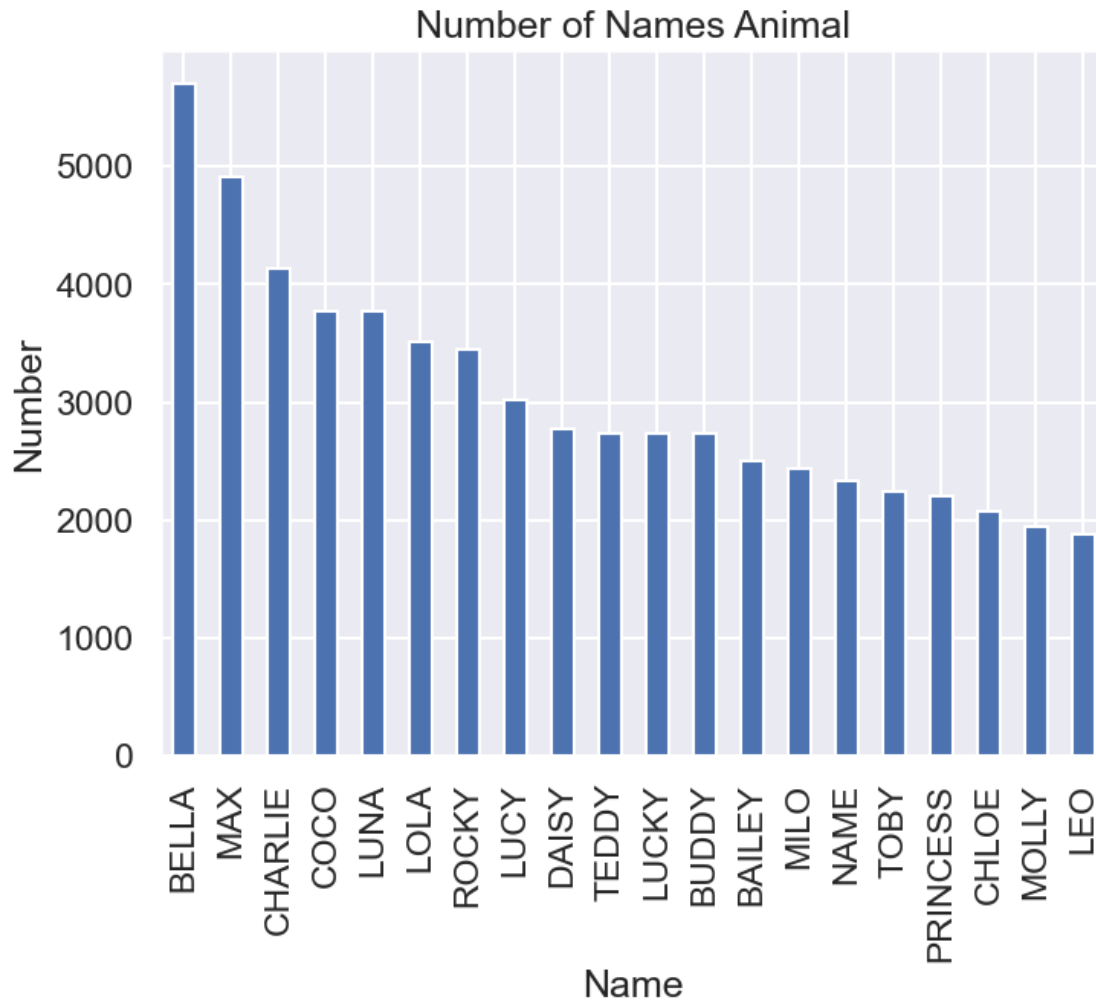
We need to clean up the data. The most popular name is unknown and second most popular is name not provided. [

```
[10]: #Drop junk names

df = df[df.AnimalName != 'UNKNOWN']
df = df[df.AnimalName != 'NAME NOT PROVIDED']

#@ 1 Show the top 20 names after dropping junk names
names = df['AnimalName'].value_counts().head(20)
```

```
names.plot.bar()
plt.xlabel("Name")
plt.ylabel("Number")
plt.title("Number of Names Animal");
```



Let us clean up the data further.

```
[11]: # DATA PREPARATION
print(df.isnull().sum())
# na values:
# AnimalName: 1027
# AnimalGender: 21
# ZipCode: 9
# LicenseExpiredDate: 77
# while we could dropna on all values, we are only doing so on AnimalGender
```

```

# we plan to predict BreedName using AnimalGender and Extract Year, so others
↳ are not required

# predicting AnimalName (or lack of name) using ZipCode and AnimalBirthYear
↳ sounds very interesting; however,
# the feedback stated that predictions using zipcodes and dates is difficult,
↳ so I think we should explore the
# relationship between these using plots instead. Such as

# easy way to show dropped columns
df_dropped = df[df['AnimalGender'].isna()]
print(df_dropped)

# drop rows that contain na in AnimalGender column
df = df[df['AnimalGender'].notna()]
#print(df.isnull().sum())

```

```

AnimalName          1027
AnimalGender         21
AnimalBirthYear      0
BreedName            0
ZipCode              9
LicenseIssuedDate    0
LicenseExpiredDate   44
Extract Year         0
dtype: int64

```

	AnimalName	AnimalGender	AnimalBirthYear	\
7356	SPARKEY	NaN	2005	
16436	SHEBAH	NaN	1997	
16684	SHEBAH	NaN	1997	
17929	SHEBAH	NaN	1997	
18002	SHEBAH	NaN	1997	
47593	BUDDY	NaN	2005	
57191	NANUK	NaN	2000	
69701	SIDNEY	NaN	2003	
110361	CHERRY	NaN	2002	
112674	NANUK	NaN	2000	
115288	SPARKEY	NaN	2005	
123852	BUDDY	NaN	2005	
142049	SIDNEY	NaN	2003	
194786	SPARKEY	NaN	2005	
198415	BUDDY	NaN	2005	
210919	NANUK	NaN	2000	
219012	SIDNEY	NaN	2003	
261017	CHERRY	NaN	2002	
261018	CHERRY	NaN	2002	
261461	BUDDY	NaN	2005	

439372      CHERRY      NaN      2002

	BreedName	ZipCode	LicenseIssuedDate	\
7356	Yorkshire Terrier	11220.0	03/16/2015	
16436	German Shepherd Dog	11218.0	07/10/2016	
16684	German Shepherd Dog	11218.0	08/09/2017	
17929	German Shepherd Dog	11218.0	10/01/2018	
18002	German Shepherd Dog	11218.0	07/31/2019	
47593	Cavalier King Charles Spaniel	10307.0	03/06/2016	
57191	Maltese	11358.0	04/19/2016	
69701	Shih Tzu	11358.0	06/10/2016	
110361	Unknown	11209.0	12/17/2016	
112674	Maltese	11358.0	04/19/2016	
115288	Yorkshire Terrier	11220.0	01/16/2017	
123852	Cavalier King Charles Spaniel	10307.0	03/08/2017	
142049	Shih Tzu	11358.0	06/06/2017	
194786	Yorkshire Terrier	11220.0	02/13/2018	
198415	Cavalier King Charles Spaniel	10307.0	03/03/2018	
210919	Maltese	11358.0	05/03/2018	
219012	Shih Tzu	11358.0	06/11/2018	
261017	Unknown	11209.0	03/02/2019	
261018	Unknown	11209.0	12/24/2020	
261461	Cavalier King Charles Spaniel	10307.0	04/02/2019	
439372	MINI PINSCHER	11209.0	12/24/2020	

	LicenseExpiredDate	Extract	Year
7356	03/01/2017		2016
16436	08/30/2017		2016
16684	08/30/2018		2017
17929	08/30/2019		2018
18002	08/30/2020		2022
47593	04/30/2017		2016
57191	03/29/2017		2016
69701	07/30/2017		2016
110361	01/30/2019		2016
112674	03/29/2018		2017
115288	03/01/2018		2017
123852	04/30/2018		2017
142049	07/30/2018		2017
194786	03/01/2019		2018
198415	04/30/2019		2018
210919	03/29/2019		2018
219012	07/30/2019		2018
261017	01/30/2021		2022
261018	01/30/2022		2022
261461	04/30/2020		2022
439372	01/30/2022		2022



Experimenting with data.

```
[12]: #Label encoding from

def Encoder(df):
    columnsToEncode = list(df.
↪select_dtypes(include=['category','object']))
    le = LabelEncoder()
    for feature in columnsToEncode:
        try:
            df[feature] = le.fit_transform(df[feature])
        except:
            print('Error encoding '+feature)
    return df
```

Machine Learning and Predictions,

```
[13]: #preserve index values

df = df.drop('ZipCode', axis=1)

df['oldIndex'] = df.index
df['oldName'] = df.index

dfConvert = df.copy()

#get rid of strings
dfE = Encoder(df)

df = pd.read_csv("http://dl.dropboxusercontent.com/s/0uzku9gb353n3nv/
↪NYC_Dog_Licensing_Dataset.csv")

df = df[df.AnimalName != 'UNKNOWN']
df = df[df.AnimalName != 'NAME NOT PROVIDED']

df = df.drop('ZipCode', axis=1)

#show top 10
dfE.head(10)

X0 = dfE.iloc[:,[0] ].values #AnimalName
X1 = dfE.iloc[:,[0,2] ].values #AnimalName, AnimalBirthYear
X2 = dfE.iloc[:,[0,2,7] ].values #AnimalName, AnimalBirthYear, Extract Year
X3 = dfE.iloc[:,[0,2,7] ].values #AnimalName, AnimalBirthYear, Extract year
X4 = dfE.iloc[:,[0,2,3] ].values #AnimalName, AnimalBirthYear, BreedName
X5 = dfE.iloc[:,[0,7] ].values #AnimalName, AnimalBirthYear
X6 = dfE.iloc[:,[0,7,3] ].values #AnimalName, Extract Year, BreedName
```

```

X7 = dfE.iloc[:,[0,7,3] ].values #AnimalName, Extract Year, BreedName
X8 = dfE.iloc[:,[0,7] ].values #AnimalName, Extract Year
X9 = dfE.iloc[:,[0,2,4,7] ].values #AnimalName, AnimalBirthYear, BreedName,
↳Extract Year
X10 = dfE.iloc[:,[0,2,4,7] ].values #AnimalName, AnimalBirthYear, BreedName,
↳Extract Year

stringX = ["1","1","1","1","1","1","1","1","1","1","1"]
stringX[0]= "AnimalName"
stringX[1]="AnimalName, AnimalBirthYear"
stringX[2]="AnimalName, AnimalBirthYear, Extract Year"
stringX[3]="AnimalName, AnimalBirthYear, Extract Year"
stringX[4]="AnimalName, AnimalBirthYear, BreedName"
stringX[5]="AnimalName, AnimalBirthYear"
stringX[6]="AnimalName, Extract Year, BreedName"
stringX[7]="AnimalName, Extract Year, BreedName"
stringX[8]="AnimalName, Extract Year"
stringX[9]="AnimalName, AnimalBirthYear, BreedName, Extract Year"
stringX[10]="AnimalName, AnimalBirthYear, BreedName, Extract Year"

#y is animal gender
y = dfE.iloc[:,[1]].values.ravel()

#y2 is breed for curiosity
y2 = dfE.iloc[:,[2]].values.ravel()

```

```
[14]: dfE.head(10)
```

```

[14]:   AnimalName  AnimalGender  AnimalBirthYear  BreedName  LicenseIssuedDate  \
0         19476             0           2014          44           1780
1         28351             1           2010          196           1780
2           462             1           2014          120           1780
3         21233             0           2013           32           1780
4         15038             0           2009          777           1780
5         11448             1           2006         1316           1780
6          3683             1           2008         1316           1780
7          4986             0           2012          697           1780
8         10884             0           2007          414           1788
9         16253             1           2009          217           1788

      LicenseExpiredDate  Extract  Year  oldIndex  oldName
0              2936      2016      0          0
1              3166      2016      1          1
2              2938      2016      2          2
3              2936      2016      3          3

```

4	3243	2016	4	4
5	3486	2016	5	5
6	3370	2016	6	6
7	3157	2016	7	7
8	1227	2016	8	8
9	2992	2016	9	9

```
[15]: dfConvert.head(10)
```

```
[15]:  AnimalName AnimalGender AnimalBirthYear \
0      PAIGE             F          2014
1      YOGI              M          2010
2       ALI              M          2014
3    QUEEN              F          2013
4     LOLA              F          2009
5      IAN              M          2006
6    BUDDY              M          2008
7 CHEWBACCA             F          2012
8  HEIDI-BO             F          2007
9   MASSIMO             M          2009
```

	BreedName	LicenseIssuedDate	LicenseExpiredDate	\
0	American Pit Bull Mix / Pit Bull Mix	09/12/2014	09/12/2017	
1	Boxer	09/12/2014	10/02/2017	
2	Basenji	09/12/2014	09/12/2019	
3	Akita Crossbreed	09/12/2014	09/12/2017	
4	Maltese	09/12/2014	10/09/2017	
5	Unknown	09/12/2014	10/30/2019	
6	Unknown	09/12/2014	10/20/2017	
7	Labrador Retriever Crossbreed	09/12/2014	10/01/2019	
8	Dachshund Smooth Coat	09/13/2014	04/16/2017	
9	Bull Dog, French	09/13/2014	09/17/2017	

	Extract	Year	oldIndex	oldName
0		2016	0	0
1		2016	1	1
2		2016	2	2
3		2016	3	3
4		2016	4	4
5		2016	5	5
6		2016	6	6
7		2016	7	7
8		2016	8	8
9		2016	9	9

```
[16]: dfBaby.head(10)
```

```
[16]:
```

	Name	Gender	Count	Probability
0	James	M	5304407	0.014517
1	John	M	5260831	0.014398
2	Robert	M	4970386	0.013603
3	Michael	M	4579950	0.012534
4	William	M	4226608	0.011567
5	Mary	F	4169663	0.011411
6	David	M	3787547	0.010366
7	Joseph	M	2695970	0.007378
8	Richard	M	2638187	0.007220
9	Charles	M	2433540	0.006660

```
[17]: dfE.index.value_counts().sum()
```

```
[17]: 490848
```

```
[18]: dfBaby['Name'] = dfBaby['Name'].str.upper()
```

Convert to upper for baby names

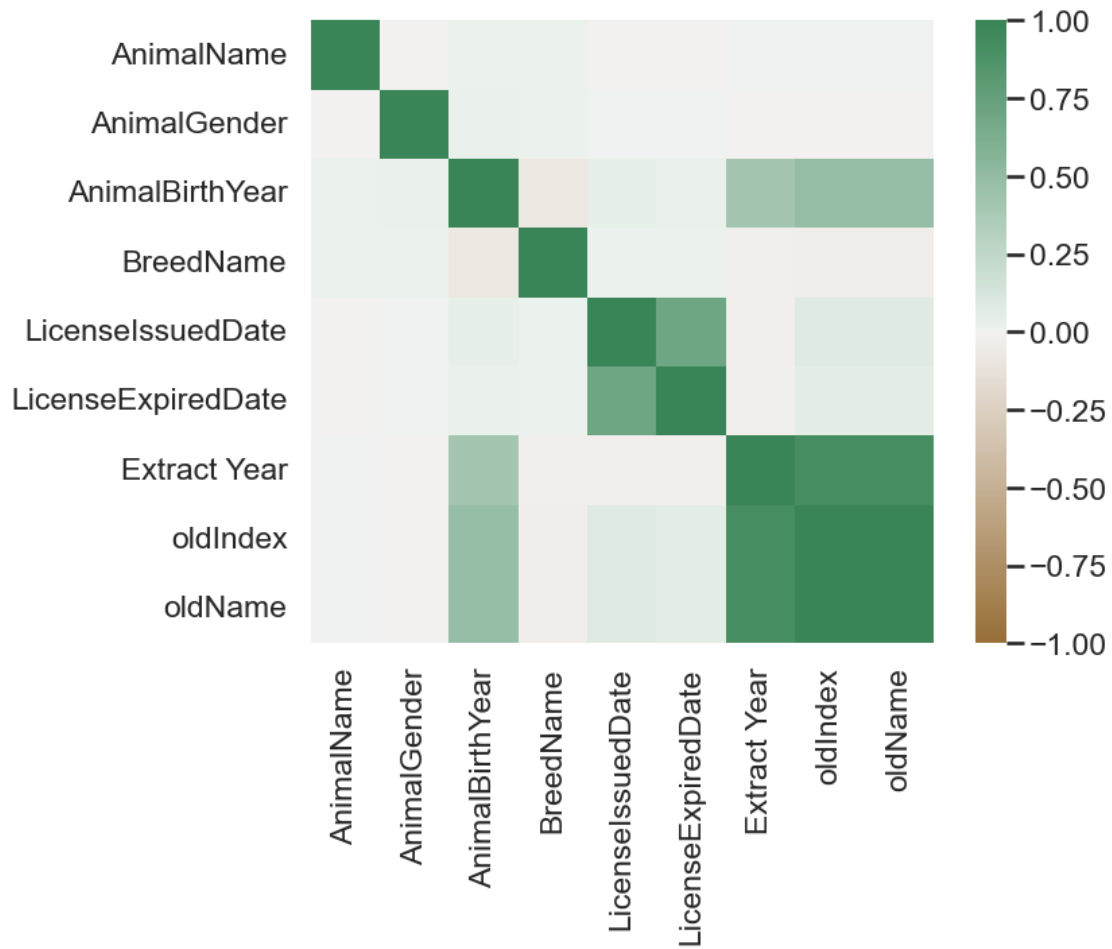
```
[19]: import matplotlib.pyplot as plt
import seaborn as sns

correlation_full_health = dfE.corr()

#cite https://seaborn.pydata.org/generated/seaborn.heatmap.html

axis_corr = sns.heatmap(
    correlation_full_health,
    vmin=-1, vmax=1, center=0,
    cmap=sns.diverging_palette(50, 500, n=500),
    square=True
)

plt.show()
```



```
[20]: dfE.head(10)
```

```
[20]:
```

	AnimalName	AnimalGender	AnimalBirthYear	BreedName	LicenseIssuedDate	\
0	19476	0	2014	44	1780	
1	28351	1	2010	196	1780	
2	462	1	2014	120	1780	
3	21233	0	2013	32	1780	
4	15038	0	2009	777	1780	
5	11448	1	2006	1316	1780	
6	3683	1	2008	1316	1780	
7	4986	0	2012	697	1780	
8	10884	0	2007	414	1788	
9	16253	1	2009	217	1788	

	LicenseExpiredDate	Extract Year	oldIndex	oldName
0	2936	2016	0	0
1	3166	2016	1	1

2	2938	2016	2	2
3	2936	2016	3	3
4	3243	2016	4	4
5	3486	2016	5	5
6	3370	2016	6	6
7	3157	2016	7	7
8	1227	2016	8	8
9	2992	2016	9	9

```
[21]: xTotal = [X0,X1,X2,X3,X4,X5,X6,X7,X8,X9,X10]
kTotal = [0 ,1 ,2 ,3 ,4 ,5 ,6 ,7 ,8 ,9 ,10]
neighborTotal = [0 , 0,0 ,0 ,0 ,0 ,0 ,0 ,0 ,0 ,0]

count = 0

results = ["1","1","1","1","1","1","1","1","1","1","1"]

for X in xTotal:

    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.25)

    scaler = StandardScaler()
    scaler.fit(X_train)
    X_train = scaler.transform(X_train)
    X_test = scaler.transform(X_test)

    for noK in range(1, 10):

        classifier = KNeighborsClassifier(n_neighbors = noK)
        classifier.fit(X_train, y_train)

        y_pred = classifier.predict(X_test)

        accuracy = accuracy_score(y_test,y_pred)

        if accuracy>neighborTotal[count]:
            neighborTotal[count]= accuracy

            kTotal[count] = noK

    results[count]= "Optimal Accuracy Score of_
↪"+str(neighborTotal[count])+ "Best K (" +str(kTotal[count])+")_
↪"+str(stringX[count])
    count += 1
```

```
for result in results:
    print(result)
```

```
Optimal Accuracy Score of 0.8678042897190169Best K (5) AnimalName
Optimal Accuracy Score of 0.8675516656887672Best K (1) AnimalName,
AnimalBirthYear
Optimal Accuracy Score of 0.6420073016493905Best K (1) AnimalName,
AnimalBirthYear, Extract Year
Optimal Accuracy Score of 0.6423577156268335Best K (1) AnimalName,
AnimalBirthYear, Extract Year
Optimal Accuracy Score of 0.856036899406741Best K (1) AnimalName,
AnimalBirthYear, BreedName
Optimal Accuracy Score of 0.7034030901623313Best K (1) AnimalName,
AnimalBirthYear
Optimal Accuracy Score of 0.6287323163178825Best K (1) AnimalName, Extract Year,
BreedName
Optimal Accuracy Score of 0.6301095247408567Best K (1) AnimalName, Extract Year,
BreedName
Optimal Accuracy Score of 0.7047314036117087Best K (1) AnimalName, Extract Year
Optimal Accuracy Score of 0.5544771497490057Best K (7) AnimalName,
AnimalBirthYear, BreedName, Extract Year
Optimal Accuracy Score of 0.5546890279679249Best K (9) AnimalName,
AnimalBirthYear, BreedName, Extract Year
```

```
[22]: xTotal = [X0,X1,X2,X3,X4,X5,X6,X7,X8,X9,X10]

count = 0

for X in xTotal:

    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.25)

    scaler = StandardScaler()
    scaler.fit(X_train)
    X_train = scaler.transform(X_train)
    X_test = scaler.transform(X_test)

    classifier = KNeighborsClassifier(n_neighbors = kTotal[count])
    classifier.fit(X_train, y_train)

    y_pred = classifier.predict(X_test)

    accuracy = accuracy_score(y_test,y_pred)

    print("Optimized Accuracy Score of "+str(count) + ": " +str(accuracy)+"\n")
    ↪("+stringX[count]+")")
```

```
count+=1
```

```
Optimized Accuracy Score of 0: 0.8636726644500945 (AnimalName)
Optimized Accuracy Score of 1: 0.8678450355303475 (AnimalName, AnimalBirthYear)
Optimized Accuracy Score of 2: 0.643963100593259 (AnimalName, AnimalBirthYear,
Extract Year)
Optimized Accuracy Score of 3: 0.6408990155811982 (AnimalName, AnimalBirthYear,
Extract Year)
Optimized Accuracy Score of 4: 0.8538366255948888 (AnimalName, AnimalBirthYear,
BreedName)
Optimized Accuracy Score of 5: 0.7048128952343699 (AnimalName, AnimalBirthYear)
Optimized Accuracy Score of 6: 0.6294575917595672 (AnimalName, Extract Year,
BreedName)
Optimized Accuracy Score of 7: 0.6326194667188213 (AnimalName, Extract Year,
BreedName)
Optimized Accuracy Score of 8: 0.7046743594758459 (AnimalName, Extract Year)
Optimized Accuracy Score of 9: 0.5543956581263446 (AnimalName, AnimalBirthYear,
BreedName, Extract Year)
Optimized Accuracy Score of 10: 0.5537437251450551 (AnimalName, AnimalBirthYear,
BreedName, Extract Year)
```

After finding the optimal predictors, we produce an accuracy score and classification report.

```
[23]: X = X1

#Animal name animal birth year

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.25)

scaler = StandardScaler()
scaler.fit(X_train)
X_train = scaler.transform(X_train)
X_test = scaler.transform(X_test)

classifier = KNeighborsClassifier(n_neighbors = 1)
classifier.fit(X_train, y_train)

y_pred = classifier.predict(X_test)

print("accuracy_score:")
accuracy = accuracy_score(y_test,y_pred)
print(accuracy)
print(" ")
print("classification_report:")
classification = classification_report(y_test, y_pred)
print(classification)
```



```
accuracy_score:  
0.8663374405111155
```

```
classification_report:  
              precision    recall  f1-score   support  
  
    0           0.86       0.85       0.85     56143  
    1           0.88       0.88       0.88     66569  
  
   accuracy                   0.87     122712  
  macro avg           0.87       0.87       0.87     122712  
weighted avg           0.87       0.87       0.87     122712
```

```
[24]: df.head(10)  
      dfE.head(10)
```

```
[24]:   AnimalName  AnimalGender  AnimalBirthYear  BreedName  LicenseIssuedDate  \  
0         19476             0           2014          44             1780  
1         28351             1           2010          196             1780  
2          462             1           2014          120             1780  
3        21233             0           2013           32             1780  
4        15038             0           2009          777             1780  
5        11448             1           2006         1316             1780  
6          3683             1           2008         1316             1780  
7          4986             0           2012          697             1780  
8        10884             0           2007          414             1788  
9        16253             1           2009          217             1788
```

```
   LicenseExpiredDate  Extract  Year  oldIndex  oldName  
0             2936       2016      0         0  
1             3166       2016      1         1  
2             2938       2016      2         2  
3             2936       2016      3         3  
4             3243       2016      4         4  
5             3486       2016      5         5  
6             3370       2016      6         6  
7             3157       2016      7         7  
8             1227       2016      8         8  
9             2992       2016      9         9
```

```
[25]: #get rid of strings  
df = Encoder(df)  
  
#drop infinities  
df = df[np.isfinite(df).all(1)]
```

```

predictors = ['AnimalName', 'AnimalBirthYear']

target = 'AnimalGender'
X = df[predictors].values
y = df[target].values

#use random state 42 70/30 split
X_train, X_test, y_train, y_test=train_test_split(X,y,test_size=.
↪30,random_state=42)

#fit the linear model
regr=LinearRegression()
regr.fit(X_train, y_train)
#get prediction
predicted=regr.predict(X_test)

rmse=np.sqrt(((y_test - predicted)**2).mean())
print('RMSE 1: {:.2f}'.format(rmse))

length_of_first = len(df['AnimalName'])
#df['BreedName'].values.reshape((-1, length_of_first))
predicted.resize((length_of_first))

sns.regplot(x=predicted, y='AnimalGender', data=df)
plt.xticks(np.arange(0,20000,500))
plt.show()

df.info()
df.head(10)

```

RMSE 1: 0.50



```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 490869 entries, 0 to 508195
Data columns (total 7 columns):
#   Column                Non-Null Count  Dtype
---  -
0   AnimalName            490869 non-null int64
1   AnimalGender          490869 non-null int64
2   AnimalBirthYear       490869 non-null int64
3   BreedName             490869 non-null int64
4   LicenseIssuedDate     490869 non-null int64
5   LicenseExpiredDate   490869 non-null int64
6   Extract Year         490869 non-null int64
dtypes: int64(7)
memory usage: 30.0 MB
```

```
[25]:   AnimalName  AnimalGender  AnimalBirthYear  BreedName  LicenseIssuedDate  \
0      19476           0           2014         44           1780
1      28351           1           2010        196           1780
2         462           1           2014        120           1780
3      21233           0           2013         32           1780
4      15038           0           2009        777           1780
5      11448           1           2006       1316           1780
6       3683           1           2008       1316           1780
```

7	4986	0	2012	697	1780
8	10884	0	2007	414	1788
9	16253	1	2009	217	1788

	LicenseExpiredDate	Extract	Year
0	2936		2016
1	3166		2016
2	2938		2016
3	2936		2016
4	3243		2016
5	3486		2016
6	3370		2016
7	3157		2016
8	1227		2016
9	2992		2016

We'll have to figure out what the next steps are given that we have a horizontal line for our fit.

Code that makes a prediction of genders based of names using a decision tree. Then its plots the actual gender count to the predicted count if you think it would be good to add here it is but if not thats okay lol the code is kind of long tho

```
[26]: import pandas as pd
from sklearn.tree import DecisionTreeClassifier
from sklearn.preprocessing import LabelEncoder
import matplotlib.pyplot as plt

# Load the data
df = pd.read_csv("http://dl.dropboxusercontent.com/s/0uzku9gb353n3nv/
↳NYC_Dog_Licensing_Dataset.csv")

# Filter out rows with NaN names and gender
df = df[['AnimalName', 'AnimalGender']].dropna()
df = df[df['AnimalGender'] != 'Unknown']

# Encode gender as numerical value
le = LabelEncoder()
df['AnimalGender'] = le.fit_transform(df['AnimalGender'])

# Split data into training and testing sets
train_size = int(len(df) * 0.8)
train_features = df['AnimalName'][:train_size]
train_labels = df['AnimalGender'][:train_size]
test_features = df['AnimalName'][train_size:]
test_labels = df['AnimalGender'][train_size:]

# Transform names to feature vectors using character frequency
char_freqs = {}
```

```

for name in train_features:
    for char in name:
        if char not in char_freqs:
            char_freqs[char] = 0
        char_freqs[char] += 1
char_freqs = {char: idx for idx, char in enumerate(sorted(char_freqs.keys()))}

train_feature_vectors = []
for name in train_features:
    freq_vec = [0] * len(char_freqs)
    for char in name:
        if char in char_freqs:
            freq_vec[char_freqs[char]] += 1
    train_feature_vectors.append(freq_vec)

test_feature_vectors = []
for name in test_features:
    freq_vec = [0] * len(char_freqs)
    for char in name:
        if char in char_freqs:
            freq_vec[char_freqs[char]] += 1
    test_feature_vectors.append(freq_vec)

# Train a decision tree classifier
clf = DecisionTreeClassifier()
clf.fit(train_feature_vectors, train_labels)

# Make predictions on the test set
predictions = clf.predict(test_feature_vectors)

# Calculate accuracy
accuracy = sum(predictions == test_labels) / len(test_labels)
print(f"Accuracy: {accuracy:.2f}")

# Predict genders for all names in the dataset
all_features = df['AnimalName']
all_feature_vectors = []
for name in all_features:
    freq_vec = [0] * len(char_freqs)
    for char in name:
        if char in char_freqs:
            freq_vec[char_freqs[char]] += 1
    all_feature_vectors.append(freq_vec)
all_predictions = clf.predict(all_feature_vectors)
predicted_genders = le.inverse_transform(all_predictions)

# Plot predicted gender distribution

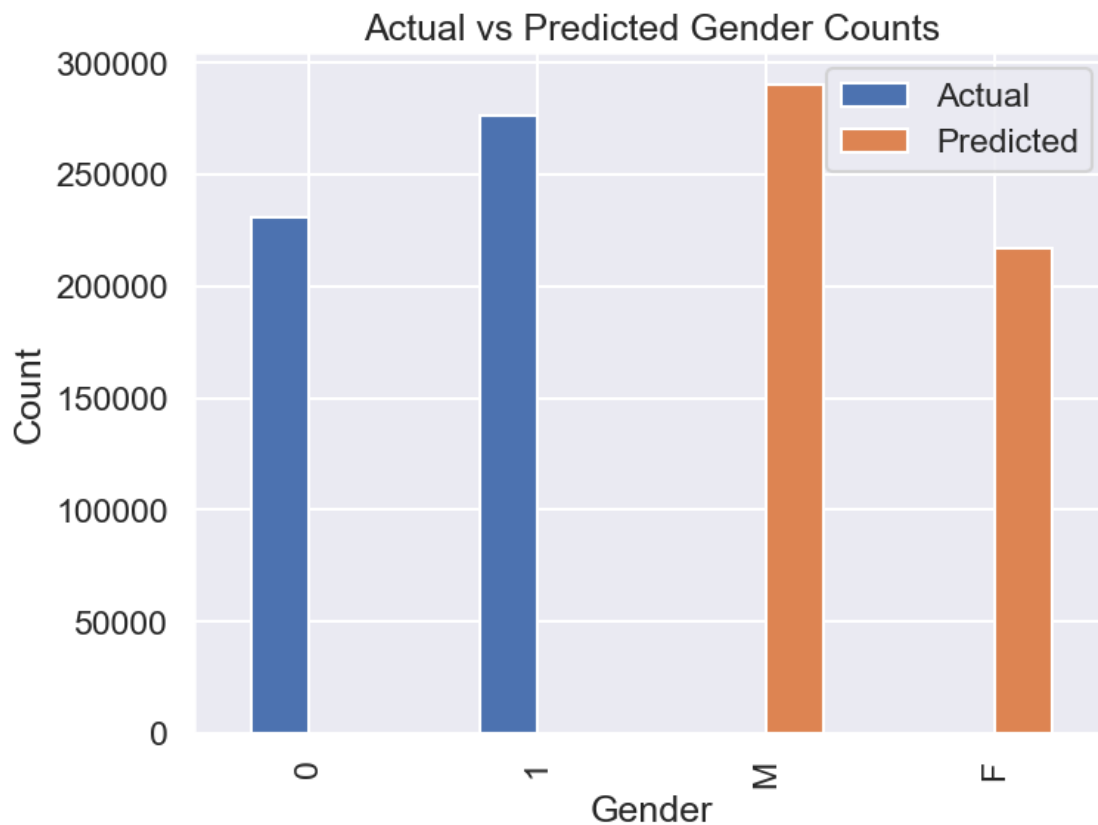
```

```

gender_counts = df.groupby('AnimalGender').size()
predicted_gender_counts = pd.Series(predicted_genders).value_counts()
combined_counts = pd.concat([gender_counts, predicted_gender_counts], axis=1)
combined_counts.columns = ['Actual', 'Predicted']
combined_counts.plot(kind='bar')
plt.title('Actual vs Predicted Gender Counts')
plt.xlabel('Gender')
plt.ylabel('Count')
plt.show()

```

Accuracy: 0.86



```

[27]: from sklearn.neighbors import KNeighborsClassifier
      from sklearn.preprocessing import LabelEncoder
      import pandas as pd

      # Filter out rows with NaN names and gender
      df = df[['AnimalName', 'AnimalGender']].dropna()
      df = df[df['AnimalGender'] != 'Unknown']

```

```

# Encode gender as numerical value
le = LabelEncoder()
df['AnimalGender'] = le.fit_transform(df['AnimalGender'])

# Transform names to feature vectors using character frequency
char_freqs = {}
for name in df['AnimalName']:
    for char in name:
        if char not in char_freqs:
            char_freqs[char] = 0
        char_freqs[char] += 1
char_freqs = {char: idx for idx, char in enumerate(sorted(char_freqs.keys()))}

feature_vectors = []
for name in df['AnimalName']:
    freq_vec = [0] * len(char_freqs)
    for char in name:
        if char in char_freqs:
            freq_vec[char_freqs[char]] += 1
    feature_vectors.append(freq_vec)

labels = df['AnimalGender']

# Split data into training and testing sets
train_size = int(len(df) * 0.8)
train_features = feature_vectors[:train_size]
train_labels = labels[:train_size]
test_features = feature_vectors[train_size:]
test_labels = labels[train_size:]

# Train and test the KNN classifier for different values of N
best_n = 0
best_accuracy = 0
for n in range(1, 21):
    clf = KNeighborsClassifier(n_neighbors=n)
    clf.fit(train_features, train_labels)
    predictions = clf.predict(test_features)
    accuracy = sum(predictions == test_labels) / len(test_labels)
    print(f"N={n}: Accuracy={accuracy:.2f}")
    if accuracy > best_accuracy:
        best_accuracy = accuracy
        best_n = n

# Print the optimal value of N
print(f"Optimal N: {best_n}")

```

```

KeyboardInterrupt                                Traceback (most recent call last)
~/var/folders/95/nfpk8m1s64791xrh28lnmvx40000gn/T/ipykernel_88010/1125926953.py
↳ in <module>
    44     clf = KNeighborsClassifier(n_neighbors=n)
    45     clf.fit(train_features, train_labels)
--> 46     predictions = clf.predict(test_features)
    47     accuracy = sum(predictions == test_labels) / len(test_labels)
    48     print(f"N={n}: Accuracy={accuracy:.2f}")

~/opt/anaconda3/lib/python3.9/site-packages/sklearn/neighbors/_classification.py
↳ in predict(self, X)
    212         Class labels for each data sample.
    213         """
--> 214         neigh_dist, neigh_ind = self.kneighbors(X)
    215         classes_ = self.classes_
    216         _y = self._y

~/opt/anaconda3/lib/python3.9/site-packages/sklearn/neighbors/_base.py in
↳ kneighbors(self, X, n_neighbors, return_distance)
    750         kwds = self.effective_metric_params_
    751
--> 752         chunked_results = list(
    753             pairwise_distances_chunked(
    754                 X,

~/opt/anaconda3/lib/python3.9/site-packages/sklearn/metrics/pairwise.py in
↳ pairwise_distances_chunked(X, Y, reduce_func, metric, n_jobs, working_memory,
↳ **kwds)
    1724         if reduce_func is not None:
    1725             chunk_size = D_chunk.shape[0]
-> 1726             D_chunk = reduce_func(D_chunk, sl.start)
    1727             _check_chunk_size(D_chunk, chunk_size)
    1728         yield D_chunk

~/opt/anaconda3/lib/python3.9/site-packages/sklearn/neighbors/_base.py in
↳ _kneighbors_reduce_func(self, dist, start, n_neighbors, return_distance)
    632         """
    633         sample_range = np.arange(dist.shape[0])[:, None]
--> 634         neigh_ind = np.argpartition(dist, n_neighbors - 1, axis=1)
    635         neigh_ind = neigh_ind[:, :n_neighbors]
    636         # argpartition doesn't guarantee sorted order, so we sort again

<__array_function__ internals> in argpartition(*args, **kwargs)

~/opt/anaconda3/lib/python3.9/site-packages/numpy/core/fromnumeric.py in
↳ argpartition(a, kth, axis, kind, order)
    837

```



```

838     """
--> 839     return _wrapfunc(a, 'argpartition', kth, axis=axis, kind=kind,
    ↪ order=order)
840
841

~/opt/anaconda3/lib/python3.9/site-packages/numpy/core/fromnumeric.py in
    ↪ _wrapfunc(obj, method, *args, **kws)
55
56     try:
---> 57         return bound(*args, **kws)
58     except TypeError:
59         # A TypeError occurs if the object does have such a method in i's

KeyboardInterrupt:

```

[ ]: